

Neuroevolution for triple-integrator control dynamics stability

Max Martinez Ruts

Abstract

Neuroevolution techniques have been successfully applied in many reinforcement learning tasks. This study aims at determining whether a neuroevolution method can be used to achieve a stable triple-integrator vehicle control system with the presence of a disturbing force. Control vehicles involving more than one integration layer are known for being very unstable and difficult to control in the presence of disturbance. Therefore a need to generate an artificial controller is originated. This study contains a description of the experiment conducted, followed by an explanation of how genetic algorithms and artificial neural networks are implemented to construct neuroevolution methods. Subsequently, the product of the experiment is discussed by analyzing the results and different experiments are executed using various parameters to analyze the effectiveness of crossover and mutation.

All code regarding this project can be found in: <https://github.com/maxmartinezruts/neuroevolution>

1 Introduction

Optimizing neural networks with genetic algorithms has proven to be a compelling approach to many control tasks such as computer game playing or robot control, as described in [3]. In such scenarios, a policy needs to be learned to determine which actions are taken by the controller. The need to use neuroevolution originates in the nature of the problem definition. The task asked is to maintain stability. However, stability cannot be instantly accomplished by a single action, but it rather depends on a sequence of actions. Meaning that a single action cannot be evaluated to be good or bad, but rather a combination of actions. These scenarios are possible to solve using reinforcement learning, a technique that consists of evaluating the effectiveness of a sequence of actions instead of singular action and awarding the solutions having greater effectiveness.

To test the effectiveness of neuroevolution a question can arise at this point to determine whether a neuroevolution method can be used to achieve a stable triple-integrator vehicle control system with the presence of a disturbing force. This study offers the employment of neuroevolution to such a task.

The main results obtained were that neuroevolution successfully achieved to generate a population with most of the vehicles being stable. Moreover, the effects crossover and mutation rate were tested to analyze their influence on the evolution process to draw conclusions for obtaining faster and stabler results.

2 Experiment description

The experiment consists of using neuroevolution (see section 4) to evolve a population composed of vehicle controllers. The control system of the vehicles is determined by an artificial neural network, which takes position p , velocity v and acceleration a as inputs, and outputs a predicted jerk input change j . The jerk input results in a change in acceleration da , at which is appended a disturbance acceleration signal at each time step to hinder the process of the jerk prediction. The following block of code pictures the vehicle control.

```
time ← 0
acc ← 0
vel ← 0
pos ← 0
while time < simulation_time do
    jerk ← PREDICT(pos, vel, acc)
    acc ← acc + jerk
    acc ← acc + random_uniform(−max_disturbance, max_disturbance)
    vel ← vel + acc
    pos ← pos + vel
    time ← time + 1
end while
```

3 Vehicle control types

Position control

Position control is the plainest type of vehicle control. A prevalent example of such is the controlling of a pointer on a computer screen with a mouse or trackpad. Nevertheless, position control is hardly applied or achieved in vehicles, with velocity and acceleration controls being more widespread.

Velocity control

Velocity control is the most frequent type of control in vehicles. It can be found in a range of different types of vehicles such as cars, and large aircraft. To control such vehicles the pilot intends to regulate the velocity of the vehicle. Thereby, the position of the vehicle is the integrated input of the control system (velocity), and acceleration is its derivative.

Acceleration control

In general, acceleration control is the hardest task a pilot will have to perform, however, it is also the most common in nature due to its parallelism to force. The most common vehicle type that is based on acceleration control is a helicopter and large aircraft, which usually operates as a double integrator system. By varying the acceleration, the velocity of the vehicle position is, therefore, the integral of the acceleration input and the position is its second integral.

Jerk control

Jerk is defined as the time derivative of acceleration. Its physical meaning is a very complex concept to understand for humans, as our vestibular system only provides us with a sense of acceleration. As one can imagine, these systems are extremely difficult to be controlled by humans. Yet, this type of control is present in some vehicles. Surface controls in aircraft, for instance, lead to small acceleration changes, meaning that the control of the surface is actually an input of the aircraft jerk. As the position or even velocity stability for these systems is extremely difficult, they are to be controlled artificially.

The scenario offered in this study is the jerk control of a vehicle.

4 Neuroevolution

Neuroevolution is the application of genetic algorithms to optimize weights and biases in a fixed topology neural network. Artificial neural networks (ANNs) are usually optimized using gradient techniques with backpropagation. Neuroevolution, alternatively, considers the use of genetic algorithms to optimize neural networks by using a reinforcement system where multiple ANNs are evaluated under a fitness score. ANNs with superior scores are reproduced more often, creating a tendency on the set of ANNs to evolve towards superior and thereby optimal ANNs.

4.1 Artificial Neural Networks

An Artificial Neural Network (ANN) is a computing system inspired by biological neural networks. Its structure is formed by connected nodes. Nodes can be seen as of biological neurons and the connections can be seen as of biological synapses, which can transmit and process signals from one node to another.

This model allows any information to be processed while traveling from one part of the ANN to another. The information is to be processed by scaling and shifting the incoming value by the use of weights and biases, and subsequently applying an activation transform to the resulting value in order to introduce non-linearities to the system. Thereby, the only values that determine how the information is to be processed are the weights and biases, which ultimately determine the behavior of the ANN.

Commonly, ANNs are structured in layers. Layered structures organize the nodes in different layers. By doing so, connections can only be present from nodes of one layer to nodes of subsequent layers, which causes the information to travel in an acyclic, unidirectional manner; from the input layer to the output layer.

The complexity of ANN goes beyond the scope of this study. From now on, therefore, an ANN will be treated as a processing function; a function processing an input to obtain an output. Treating an ANN as such is useful to envision what is the actual role of an ANN in this study. ANNs are used for building a function that can process the current information of the vehicle control to return an action as a response.

4.1.1 Neural Network Topology

The topology of an ANN refers to its structural organization of nodes and connections. Topology of ANNs can contribute to their performance, as described in [9].

Having described the scenario in section 2, the input and output layers can already be determined. The input layer consists of three nodes - position p , velocity v and acceleration a - values. The output layer consists of one node containing the information on the increase in acceleration da . An extra hidden layer consisting of 6 nodes is also present to add depth - and therefore complexity - to the neural network. In order to account for non-linearities within the input variables to predict the system's behavior, the activation function of the hidden and output layers are sigmoid functions, (used for probabilistic examples, as it outputs values from 0 to 1). The outputs are then mapped from $[0,1]$ to $[-0.5,0.5]$ to obtain an unbiased jerk input. Figure 1 displays the topology of the ANN to be used.

The neural network can, therefore, be seen as a function that inputs the position, velocity, and acceleration of the vehicle, and outputs a change in acceleration. Hence, the problem scenario consists of optimizing ANN such that it can precisely determine da to achieve stability.

There are several methods to optimize an ANN depending on the behavior and the utility of the ANN. The method used for this scenario is neuroevolution; a genetic algorithm approach to optimize reinforcement learning ANNs. The optimization of the ANN will be delivered by using neuroevolution to change the weights

and biases of the ANN. However, a more elaborated and complex behavior of an ANN could be achieved by also altering the topology of the ANN. This method is referred to as NEAT (Neuroevolution of Augmented Topologies) presented by [8].

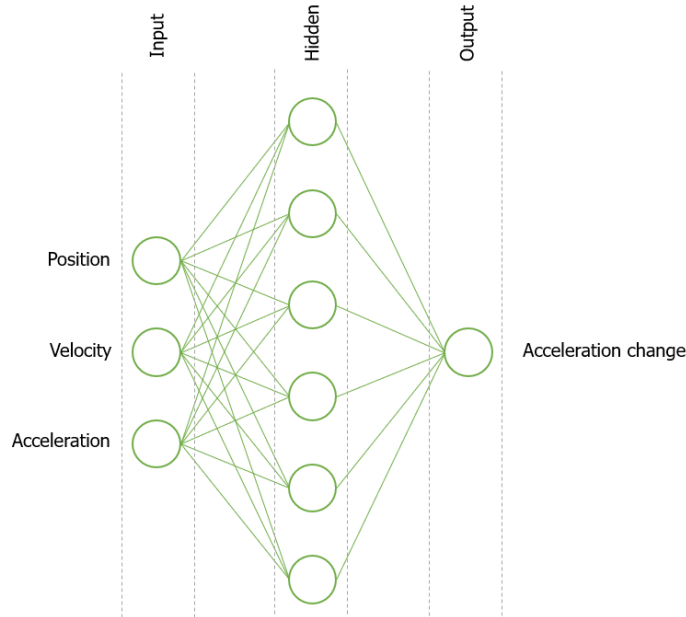


Figure 1: Neural Network Topology

4.2 Genetic Algorithms

Genetic Algorithms (GAs) are inspired by Darwins theory of biological evolution [1]. GAs differ from the rest of optimization methods by the fact that a set of solutions is maintained rather than a singular one. GAs are formed by a set of genetically diverse specimens, being evaluated under a certain fitness function. Each specimen inherits a likelihood of survival, proportional to its evaluated fitness score. This behavior is often referred to as survival of the fittest, the basic argument on which natural selection is based. This behavior allows species to evolve, as specimens which survive longer are also the ones with more chances to reproduce similar copies of themselves, thereby inducing a tendency on the population of subsequent generations towards fitter specimens.

GAs are proved to be very effective in many real world and engineering optimization tasks such as electromagnetic system design [5] and aerodynamic design [7]. The reason why it has become widely used and effective is due to its simplistic implementation in object-oriented programming. A parallelism is present between biological evolution and GAs, populations being sets, specimens being objects, genes being weights and biases and brains being ANNs. One could think of this parallelism by comparing the biological approach:

"Natural selection invokes a tendency on succeeding generations to produce populations with specimens having smarter brains as a consequence of the evolution of their DNA. Such improvement leads to a general rise on their performance"

With the programmatic approach:

"Genetic algorithms invokes a tendency on succeeding generations to produce sets with objects having optimal neural networks as a consequence of the evolution of their weights and biases. Such improvement leads to a general rise in their fitness score."

4.2.1 Selection

Selection is the determination process to decide which specimens will leave a successor on the next generation and which specimens will disappear without a successor. GAs use a different variety of selection methods. Most of them, however, have a peculiarity in common; they tend to select fit specimens, based on the behavior known as survival of the fittest.

Fitness Score determination

To determine the selection process, a fitness score has to be awarded to each specimen in the population. In the scenario proposed, a fitness score will be proportional to the stability achieved by the vehicle (specimen). There exist several methods to evaluate the stability of a vehicle. The method used consists on running a simulation of the vehicles, departing from the center of the screen and appending a score at each timestamps that is higher the closest the vehicle is from the center of the screen, meaning that vehicles remaining close to the center of the screen during the simulation (hence achieving stability) will receive a higher fitness score. The exact award given at each time step is determined using a Gaussian function:

$$d_score = e^{-(pos-center)^2} \quad (1)$$

This ensures that only vehicles close to the center of the screen will be awarded a significant fitness score (see Figure 2).




Figure 2: Gaussian Function

The following code provides a programmatic approach picturing how the fitness score is determined by executing a simulation. As stated in the experiment description, each time step an extra acceleration term is randomly introduced.

```

function EVALUATE_FITNESS(specimen)
  score  $\leftarrow$  0
  time  $\leftarrow$  0
  acc  $\leftarrow$  0
  vel  $\leftarrow$  0
  pos  $\leftarrow$  0
  while time < simulation_time do
    jerk  $\leftarrow$  PREDICT(pos, vel, acc, specimen.genotype)
    acc  $\leftarrow$  acc + jerk
    acc  $\leftarrow$  acc + random_uniform( $-max\_disturbance$ ,  $max\_disturbance$ )
    vel  $\leftarrow$  vel + acc
    pos  $\leftarrow$  pos + vel
    score  $\leftarrow$  score +  $e^{-(pos-center)^2}$ 
    time  $\leftarrow$  time + 1
  end while
  return score
end function

```

Selection method

The selection method used for this specific scenario is based on the principle of survival of the fittest. The population size is set constant to N throughout all generations. After all specimens of the population are evaluated under their fitness scores, an algorithm is applied to decide which of the specimens will partake in the reproduction process. This algorithm assigns a probability of reproduction equivalent to its fitness score, therefore leaving fitter specimens with higher chances to take part in the reproduction process.

The following block of code illustrates how a progenitor for a newly generated specimen is selected. Note that the method only works under the condition that the total score generated by all the population is equal to 0. Therefore, the fitness scores have to be previously normalized such that the sum of all scores is 1:

```

function PICK_PARENT
  i  $\leftarrow$  0
  r  $\leftarrow$  random(0, 1)
  while r > 0 do
    r  $\leftarrow$  r - EVALUATE_FITNESS(specimens[i])
    i  $\leftarrow$  i + 1
  end while
  i  $\leftarrow$  i - 1
  return specimens[i]
end function

```

4.2.2 Reproduction

In nature, reproduction is the process of generating new specimens sharing similar genetic information to the one of their progenitors. Multiple reproduction methods are present in nature. However, they all are designed to ensure diversity of a species. Genetic diversity is achieved by DNA alternation, which can be developed by many different processes. Two of the most common are crossover and mutation.

Crossover

Crossover is the ability for some specimens to reproduce a hybridized genetic copy of themselves. Crossover is, therefore, a feature that allows the species to diversify. Crossover is commonly thought as sexual reproduction, which involves the union of a male and a female specimen. Although nature mostly experiences sexual reproduction with two progenitors, a crossover method can be designed where multiple progenitors are

involved in the contribution of genes to the successor. However, in the experiment, bi-parent reproduction is chosen. A programmatic approach to deal with crossover is to simply create a hybrid genotype from the predecessors.

As genotypes are composed by the weights and biases of the ANN, which are arranged as matrices, a hybrid genotype can be generated by simply creating a weighted average of the individual genotypes of each progenitor. The following block of code pictures a programmatic approach to achieve crossover.

```
function CROSSOVER
  genotype1 ← PICK_PARENT
  genotype2 ← PICK_PARENT
  genotype ← (genotype1 + genotype2)/2
  return genotype
end function
```

As the genetic information needed to model the ANN is composed by four different genotypes (*weights_hidden*, *biases_hidden*, *weights_output*, *biases_output*), the crossover method will be applied to each of these four genotypes.

Mutation

Mutation is the capability for a specimen of modifying its genotype, providing the population the ability to diversify. Diversity is key-driving in GAs, as it allows a population to have a broad genotype domain, which expands the search domain, thereby increasing the possibilities from emerging new species with a unique set of genes that drive specimens to increase their performance.

If diversification was nonexistent, a population would tend to converge to a single species. Certainly, it would be the best species generated at the moment, but the new species would not tend to evolve, as all individuals would tend to the fittest specimen in the population without the opportunity to generate specimens with distinct genotypes. It is therefore meaningful to define a parameter that represents the likelihood for a gene of a specimen to be cloned from its progenitor. Such a parameter is referred to as mutation rate; defined as the likelihood for a gene to be mutated.

If the mutation rate is set high, the population tends to diversify. However, the successors might not carry enough genetic information from their progenitors, and thereby some of the genes that made the progenitor a fit specimen could be lost. Similarly, if the mutation rate is set low, the progenitors will carry most of the genetic information of their successors but this could lead to a poor tendency to evolve, as the diversity of the species might be too low. Another relevant parameter is the mutation magnitude. When a gene of the specimen is mutated, it can either be redefined as a random gene or as a variation of the original gene. In the last case, a key parameter is the mutation magnitude; defined as the magnitude of how much is the original gene varied.

The following block of code uses a random mutation magnitude, which can contribute to increasing the diversity of the population, as it adds a random parameter to the method.

```
function MUTATION(genotypes, mutation_rate)
  for all gene in genotype do
    r ← random(0,1)
    if r < mutation_rate then
      gene ← gene + random_normal
    end if
  end for
end function
```

5 Backpropagation vs GAs for ANN optimization

In Neural-Network optimization, backpropagation is the most common approach. Backpropagation is a gradient search technique. Such technique is based on calculation the optimal direction in the domain space (weights and biases) such that the loss is minimized, by the use of the chain rule. The loss function is defined as the distance between the output of the neural network and the desired output. To use backpropagation it is, therefore, necessary to know the real, desired output. This limits the use of backpropagation to supervised learning, as it is not possible to optimize the neural network if the solution (output) is not known. Moreover, gradient descent techniques guarantee the best solution in the region of the starting point. Obtaining a global solution is therefore dependent on the choice of initial starting values.

GAs tend to produce global solutions (see [6]). Specimens in a global solution are fitter than specimens in a local solution, thereby inducing a tendency on the population towards global solutions. Due to this behavior, GAs typically outperform backpropagation for very non-linear problems as described in [4], as gradient descent tends to produce local solutions due to the loss function being hilly among the search space. In GAs, instead, if the population size is kept high and the diversity of the population is high enough such that a broad space can be covered, a global solution is encountered. Another essential advantage over backpropagation is that GAs can be used for unsupervised learning and reinforcement learning as described in [2], as it is not based on a loss function (dependent on the desired output) but on a fitness function (dependent on the performance achieved by the input).

The previous reasoning leaves GAs to be the best fit for the problem scenario. The main two reasons are that the solution to the problem requires reinforcement learning and that the solution is rather non-linear. But why is the case that this scenario requires reinforcement learning?

The proposed problem lies in a scenario where the neural network has as inputs (p,v and a) and output (da). However, given the inputs, the best output cannot be directly determined, as it is difficult to determine if the action will produce a good or a bad result in order to achieve stability. In other words, a loss function cannot be determined. This is because an acceleration change will not produce a direct result in the stability of the object. The stability will be rather achieved by a sequence of good acceleration changes. Therefore singular acceleration change wont have a direct result on the stability of the object, but rather a chaotic one. It is for this reason why an output cannot be directly compared to the desired output, precluding the determination of a loss function and therefore precluding supervised learning. A similar situation is encountered in chess. One can not train a supervised neural network by determining which is the best move to do in a certain situation, as the effectuation of a move will not have an instant result on the game but rather a chaotic one that will affect the result of the game in the future.

6 Results & Discussion

This section provides a close look at the results obtained by executing several experiments using neuroevolution.

6.1 Experiment setup

In the following sections, the experiment will be executed several times to compare the effects of different parameters. Unless specifically stated, the parameters used for the experiment executions are the following:

- Input layer: 3 nodes
- Hidden layer: 6 nodes, using sigmoid activation function
- Output layer: 3 nodes, using sigmoid activation function
- Population size: 50 vehicles
- Number of generations: 50
- Maximum acceleration disturbance: $0.05 [m/s^2]$
- Simulation time: 500 time steps
- Reproduction methods: Crossover + Mutation
- Mutation rate: 10 %
- Mutation magnitude: random number generated from normal random distribution

6.2 Results analysis

The following section presents the discussion of the results obtained by executing a simulation with the environment setup presented in subsection 6.1. The following progression graph (see Figure 3) contains the scores of the best and worst performing vehicles for each generation as well as the median and mean scores of the entire population.

The results are promising; applying neuroevolution to a set of only 50 vehicles during 50 generations results in a population having most of the vehicles capable of achieving a stable jerk control. Achieving such implies that the ANN is capable of returning the right jerk value to achieve a stable position (third integral from jerk) with the extra difficulty of the presence of an unpredictable acceleration disturbance.

By observing Figure 3, it can be observed that the first few generations show a slow evolution progression. This could be mainly due to the low mutation rate used, leading to a poor diversification of the population. As commented before, low diversification hinders the appearance of new species. The population seems to stabilize at generation 10, by almost maintaining the same medians and the same means during 20 more generations. However, at generation 30, something remarkable occurs; a new specimen with a successful genotype is created, and due to the low mutation rate used subsequent populations rapidly acquire similar genotypes.

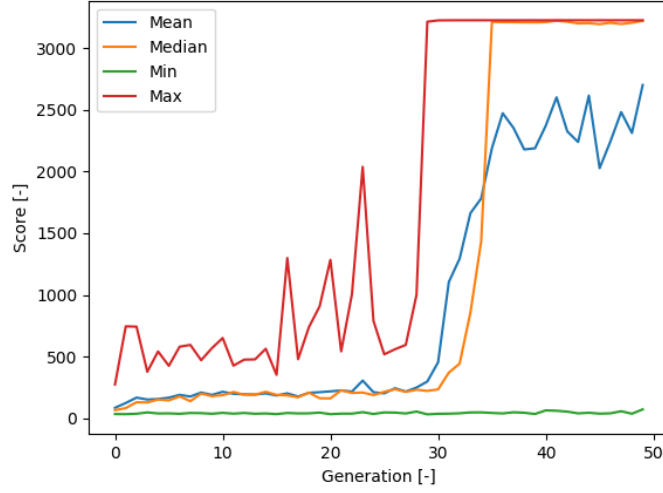


Figure 3: Evolution of scores over 50 generations

Notice that using a low mutation rate causes the population to be closely related to previous generations. This behavior can be observed by comparing 'max' and 'median' progressions. A slight shift of around 7 generations can be observed between the two progression lines (see Figure 3). After one specimen achieves a high fitness score, the low mutation rate used enables the following generations to rapidly generate more specimens with similar genotypes, thereby plaguing the entire population by similar successful specimens. See for example generation 30 (see Figure 4). A perfect-score specimen appears. By generation 37 (see Figure 5), most of the population has already acquired a similar genotype to the one of the successful specimen encountered at generation 30.




Figure 4: Histogram of scores at generation 30




Figure 5: Histogram of scores at generation 37

6.3 Effect of crossover

Crossover is very common in nature, as it is present in all sexual reproduction species. To investigate the effect of crossover for this scenario, two identical experiments are executed; one including bi-parent crossover and one including mono-parent reproduction. The mono-parent reproduction simply consists of copying the genetic information of one parent (see programmatic approach).

```
function MONOPARENT
    genotype  $\leftarrow$  PICK_PARENT
    return genotype
end function
```

Comparing both experiments (see Figure 7 and Figure 6), it can be observed that both methods achieve a successful evolution. However, their progression behavior is different. The experiment using mono-parent reproduction seems to follow a linear progression, whereas the crossover experiment seems to rapidly evolve during the first generations and stabilize afterward. This difference is present due to the effect on the diversity imposed by the crossover method. Since the resulting specimen contains hybridized information from two different progenitors, a completely new specimen is generated, thereby increasing diversity.

As commented before, diversity does not imply successful evolution, but it rather implies a behavior of evolution. In this scenario, it can be seen that crossover would be a more efficient method of reproduction during the first few generations. However, mono-parenting would be a better option for subsequent generations. The reason why diversity is important during the first generations is that there is a need to generate different specimens to have a broad search domain to increase the likelihood to find fit specimens. However, once the population is composed mainly by fit specimens, diversification will difficult the evolution process, as more genetic information from progenitors will be lost.




Figure 6: Evolution of scores over 50 generations using mono-parental reproduction




Figure 7: Evolution of scores over 50 generations using crossover reproduction

6.4 Effect of mutation rate

As commented before, the mutation rate is the parameter that determines how often mutation occurs. To investigate the effect of mutation rate, three identical experiments are executed with distinct mutation rates. The first experiment is provided with a mutation rate of 10 %, the second one with a mutation rate of 1 % and the third one with a mutation rate of 0.1 %. To dramatize the effect of the mutation rate, the value of

the gene is determined differently; when a mutation occurs, the random value is not added to the gene but rather substituted by the gene. By doing so, a mutation results in a more drastic effect on the genotype, as no previous information of the gene is kept.

The first experiment, executed with a mutation rate of 10 % results in a high diversification of the species. As commented in subsubsection 4.2.2, a high diversification enables a population to generate very spread genotypes, thereby increasing the probability of discovering good performing specimens. However, this behavior results in poor genetic information conservation from the progenitors, resulting in a poor tendency to evolve. As it can be observed in (see Figure 8), the population does not struggle to generate perfect-score specimens; however, the genetic information of these specimens is frequently not carried to the next generation and therefore is lost, resulting in a poor evolution over the generations.




Figure 8: Evolution of scores over 50 generations using 10 % Mutation Rate

The second experiment, executed with a mutation rate of 1 % gives much better results. If diversification is reduced, specimens can carry more genetic information from their progenitors, resulting in a stronger tendency to evolve towards fitter populations. This behavior can be observed in generations 22 and 43 (see Figure 9). In these generations, high-performance specimens are generated, and the succeeding generations rapidly adopt their genetic information. This behavior results in a rich evolution, having a good balance between diversity and ability to conserve genetic information.

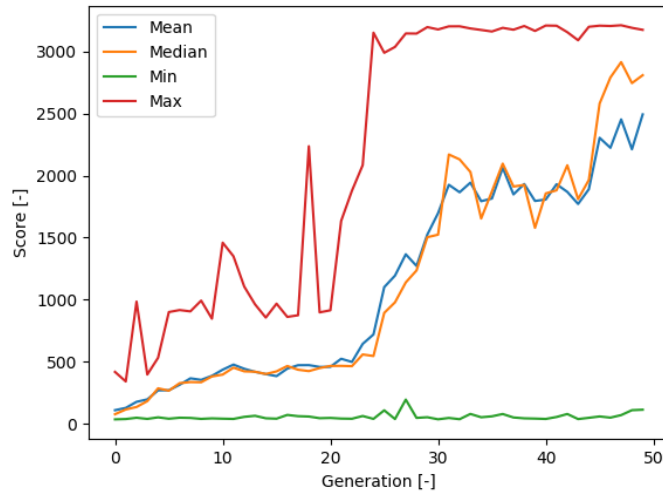


Figure 9: Evolution of scores over 50 generations using 1 % Mutation Rate

The third experiment is executed using a mutation rate of 0.1%. This rate results in a practically nonexistent mutation. It is important to be aware that the only reason why the population seems to evolve (see Figure 10) is due to the diversification imposed by the crossover process. However, the mutation factor barely has any influence on the diversification of the population. Comparing this results with the previous experiment, it can be observed that, although the specimens can carry most of the genetic information from their parents, their ability to diversify is too poor to achieve a collective evolution.




Figure 10: Evolution of scores over 50 generations using 0.1 % Mutation Rate

References

- [1] D. E. Goldberg. Genetic algorithms in search, optimization and machine learning. *Reading, MA: Addison-Wesley.*, 1989.
- [2] F. Gomez and R. Miikkulainen. Learning robust nonlinear control with neuroevolution. *Department of Computer Sciences, The University of Texas at Austin*, 2001.
- [3] F. Gomez and R. Miikkulainen. Solving non-markovian control tasks with neuroevolution. *Department of Computer Sciences, The University of Texas at Austin*, 2001.
- [4] F. Gomez and R. Miikkulainen. Learning robust nonlinear control with neuroevolution. *Department of Computer Sciences, The University of Texas at Austin*, 2001.
- [5] E. Michielssen and D. Weile. Electromagnetic system design using genetic algorithms. *JohnWiley Sons Ltd*, 1995.
- [6] M. Niknafs. Neural network optimization. 2008.
- [7] J. Periaux. Robust genetic algorithms for optimization problems in aerodynamic design. *John Wiley Sons Ltd*, 1995.
- [8] K. O. Stanley and R. Miikkulainen. Evolving neural networks through augmenting topologies. *The MIT Press Journals*, 10(2).
- [9] J. Torres, M. Munoz, J. Marro, and P. Garrido. Influence of topology on the performance of a neural network. *Department of Electr. and Matter Physics, and Institute Carlos I for Theoretical and Computational Physics, University of Granada, E-18071 Granada, Spain*, 2004.